

# Data Understanding

The high level goal of analyzing the Aviation Accident Database & Synopses, up to 2023 Data set is to determine which aircraft has the lowest risk to ensure the company starts a new business endeavour. I'll do that by determining which models causes the most accidents per year. When the Dataset was being collected in 2023, there were six fatal accidents globally in 2023, with these resulting in 115 onboard deaths, and this triggered a need to look for a solution.

# Data Understanding

```
In [ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
```

```
In [ ]: import zipfile
```

```
In [ ]: #Reading the Dataset
df = pd.read_csv("AviationData.csv (1).zip", encoding="iso-8859-1", low_memory=False)
```

```
In [ ]: #Checking for number of rows and columns
df.shape
```

```
Out[ ]: (88889, 31)
```

```
In [ ]: df.tail()
```

Out [ ]:

	Event.Id	Investigation.Type	Accident.Number	Event.Date	Location	Coun
88884	20221227106491	Accident	ERA23LA093	2022-12-26	Annapolis, MD	Unit Sta
88885	20221227106494	Accident	ERA23LA095	2022-12-26	Hampton, NH	Unit Sta
88886	20221227106497	Accident	WPR23LA075	2022-12-26	Payson, AZ	Unit Sta
88887	20221227106498	Accident	WPR23LA076	2022-12-26	Morgan, UT	Unit Sta
88888	20221230106513	Accident	ERA23LA097	2022-12-29	Athens, GA	Unit Sta

5 rows × 31 columns

In [ ]: df.head()

Out [ ]:

	Event.Id	Investigation.Type	Accident.Number	Event.Date	Location	Countr
0	20001218X45444	Accident	SEA87LA080	1948-10-24	MOOSE CREEK, ID	Uniter State
1	20001218X45447	Accident	LAX94LA336	1962-07-19	BRIDGEPORT, CA	Uniter State
2	20061025X01555	Accident	NYC07LA005	1974-08-30	Saltville, VA	Uniter State
3	20001218X45448	Accident	LAX96LA321	1977-06-19	EUREKA, CA	Uniter State
4	20041105X01764	Accident	CHI79FA064	1979-08-02	Canton, OH	Uniter State

5 rows × 31 columns

## Checking for summary statistics

In [ ]: #Checking for summary statistics  
df.describe()

Out[ ]:

	Number.ofEngines	Total.Fatal.Injuries	Total.Serious.Injuries	Total.Minor.Injuries	Tot
count	82805.000000	77488.000000	76379.000000	76956.000000	ε
mean	1.146585	0.647855	0.279881	0.357061	
std	0.446510	5.485960	1.544084	2.235625	
min	0.000000	0.000000	0.000000	0.000000	
25%	1.000000	0.000000	0.000000	0.000000	
50%	1.000000	0.000000	0.000000	0.000000	
75%	1.000000	0.000000	0.000000	0.000000	
max	8.000000	349.000000	161.000000	380.000000	

◀

▶

## Checking for Data types

In [ ]:

```
#Checking for the Data types
df.dtypes
```

```
Out[ ]: Event.Id                object
Investigation.Type            object
Accident.Number               object
Event.Date                    object
Location                      object
Country                       object
Latitude                      object
Longitude                     object
Airport.Code                  object
Airport.Name                  object
Injury.Severity               object
Aircraft.damage               object
Aircraft.Category             object
Registration.Number           object
Make                          object
Model                         object
Amateur.Built                 object
Number.of.Engines             float64
Engine.Type                   object
FAR.Description               object
Schedule                      object
Purpose.of.flight             object
Air.carrier                   object
Total.Fatal.Injuries          float64
Total.Serious.Injuries         float64
Total.Minor.Injuries          float64
Total.Uninjured               float64
Weather.Condition              object
Broad.phase.of.flight         object
Report.Status                  object
Publication.Date               object
dtype: object
```

## Checking for missing values

```
In [ ]: #Checking for missing values
df.isna().sum()
```

```
Out[ ]: Event.Id                0
        Investigation.Type      0
        Accident.Number        0
        Event.Date             0
        Location               52
        Country                226
        Latitude               54507
        Longitude              54516
        Airport.Code           38757
        Airport.Name           36185
        Injury.Severity        1000
        Aircraft.damage         3194
        Aircraft.Category       56602
        Registration.Number     1382
        Make                   63
        Model                  92
        Amateur.Built          102
        Number.of.Engines       6084
        Engine.Type             7096
        FAR.Description         56866
        Schedule               76307
        Purpose.of.flight       6192
        Air.carrier             72241
        Total.Fatal.Injuries    11401
        Total.Serious.Injuries  12510
        Total.Minor.Injuries    11933
        Total.Uninjured        5912
        Weather.Condition       4492
        Broad.phase.of.flight   27165
        Report.Status           6384
        Publication.Date        13771
        dtype: int64
```

## Checking for Duplicates

```
In [ ]: df.duplicated().sum()
```

```
Out[ ]: np.int64(0)
```

```
In [ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 88889 entries, 0 to 88888
Data columns (total 31 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Event.Id                             88889 non-null  object
1   Investigation.Type                    88889 non-null  object
2   Accident.Number                      88889 non-null  object
3   Event.Date                           88889 non-null  object
4   Location                             88837 non-null  object
5   Country                             88663 non-null  object
6   Latitude                             34382 non-null  object
7   Longitude                            34373 non-null  object
8   Airport.Code                         50132 non-null  object
9   Airport.Name                         52704 non-null  object
10  Injury.Severity                      87889 non-null  object
11  Aircraft.damage                      85695 non-null  object
12  Aircraft.Category                    32287 non-null  object
13  Registration.Number                 87507 non-null  object
14  Make                                88826 non-null  object
15  Model                               88797 non-null  object
16  Amateur.Built                       88787 non-null  object
17  Number.of.Engines                   82805 non-null  float64
18  Engine.Type                         81793 non-null  object
19  FAR.Description                     32023 non-null  object
20  Schedule                            12582 non-null  object
21  Purpose.of.flight                   82697 non-null  object
22  Air.carrier                         16648 non-null  object
23  Total.Fatal.Injuries                 77488 non-null  float64
24  Total.Serious.Injuries               76379 non-null  float64
25  Total.Minor.Injuries                 76956 non-null  float64
26  Total.Uninjured                     82977 non-null  float64
27  Weather.Condition                   84397 non-null  object
28  Broad.phase.of.flight                61724 non-null  object
29  Report.Status                       82505 non-null  object
30  Publication.Date                     75118 non-null  object
dtypes: float64(5), object(26)
memory usage: 21.0+ MB
```

```
In [ ]: df.index
```

Out[ ]: RangeIndex(start=0, stop=88889, step=1)

```
In [ ]: df.value_counts()
```

Out[ ]: Series([], Name: count, dtype: int64)

```
In [ ]: df.values
```

```
Out[ ]: array([[ '20001218X45444', 'Accident', 'SEA87LA080', ..., 'Cruise',  
                'Probable Cause', nan],  
              [ '20001218X45447', 'Accident', 'LAX94LA336', ..., 'Unknown',  
                'Probable Cause', '19-09-1996'],  
              [ '20061025X01555', 'Accident', 'NYC07LA005', ..., 'Cruise',  
                'Probable Cause', '26-02-2007'],  
              ...,  
              [ '20221227106497', 'Accident', 'WPR23LA075', ..., nan, nan,  
                '27-12-2022'],  
              [ '20221227106498', 'Accident', 'WPR23LA076', ..., nan, nan, nan],  
              [ '20221230106513', 'Accident', 'ERA23LA097', ..., nan, nan,  
                '30-12-2022']], dtype=object)
```

## Data Cleaning

```
In [ ]: #Checking if there are duplicates  
df.duplicated().sum()
```

```
Out[ ]: np.int64(0)
```

```
In [ ]: df.isna().sum()
```

```
Out[ ]: Event.Id                0
        Investigation.Type      0
        Accident.Number        0
        Event.Date             0
        Location                52
        Country                 226
        Latitude                54507
        Longitude               54516
        Airport.Code            38757
        Airport.Name            36185
        Injury.Severity         1000
        Aircraft.damage         3194
        Aircraft.Category       56602
        Registration.Number     1382
        Make                    63
        Model                   92
        Amateur.Built           102
        Number.ofEngines        6084
        Engine.Type             7096
        FAR.Description         56866
        Schedule                76307
        Purpose.of.flight       6192
        Air.carrier             72241
        Total.Fatal.Injuries    11401
        Total.Serious.Injuries  12510
        Total.Minor.Injuries    11933
        Total.Uninjured         5912
        Weather.Condition       4492
        Broad.phase.of.flight   27165
        Report.Status           6384
        Publication.Date        13771
        dtype: int64
```

```
In [ ]: df.isna().sum()
```



```
Out[ ]: Event.Id          0
        Investigation.Type 0
        Accident.Number    0
        Event.Date         0
        Location           52
        Country            226
        Latitude           54507
        Longitude          54516
        Airport.Code       38757
        Airport.Name       36185
        Injury.Severity    1000
        Aircraft.damage    3194
        Aircraft.Category  56602
        Registration.Number 1382
        Make              63
        Model             92
        Amateur.Built     102
        Number.of.Engines  6084
        Engine.Type       7096
        FAR.Description    56866
        Schedule          76307
        Purpose.of.flight  6192
        Air.carrier       72241
        Total.Fatal.Injuries 11401
        Total.Serious.Injuries 12510
        Total.Minor.Injuries 11933
        Total.Uninjured    5912
        Weather.Condition  4492
        Broad.phase.of.flight 27165
        Report.Status      6384
        Publication.Date   13771
        dtype: int64
```

```
In [ ]: df.isnull().sum()
```

```
Out[ ]: Event.Id          0
Investigation.Type      0
Accident.Number        0
Event.Date             0
Location               52
Country                226
Latitude               54507
Longitude              54516
Airport.Code           38757
Airport.Name           36185
Injury.Severity        1000
Aircraft.damage        3194
Aircraft.Category      56602
Registration.Number    1382
Make                   63
Model                  92
Amateur.Built          102
Number.ofEngines       6084
Engine.Type            7096
FAR.Description        56866
Schedule               76307
Purpose.of.flight     6192
Air.carrier            72241
Total.Fatal.Injuries   11401
Total.Serious.Injuries 12510
Total.Minor.Injuries   11933
Total.Uninjured        5912
Weather.Condition      4492
Broad.phase.of.flight  27165
Report.Status          6384
Publication.Date       13771
dtype: int64
```

```
In [ ]: df.dropna()
```

```
Out[ ]:  Event.Id  Investigation.Type  Accident.Number  Event.Date  Location  Country  Latitude
0 rows x 31 columns
```

```
In [ ]: df.duplicated().sum()
```

```
Out[ ]: np.int64(0)
```

```
In [ ]: print(df.columns)
```

```
Index(['Event.Id', 'Investigation.Type', 'Accident.Number', 'Event.Date',
      'Location', 'Country', 'Latitude', 'Longitude', 'Airport.Code',
      'Airport.Name', 'Injury.Severity', 'Aircraft.damage',
      'Aircraft.Category', 'Registration.Number', 'Make', 'Model',
      'Amateur.Built', 'Number.of.Engines', 'Engine.Type', 'FAR.Description',
      'Schedule', 'Purpose.of.flight', 'Air.carrier', 'Total.Fatal.Injuries',
      'Total.Serious.Injuries', 'Total.Minor.Injuries', 'Total.Uninjured',
      'Weather.Condition', 'Broad.phase.of.flight', 'Report.Status',
      'Publication.Date'],
      dtype='object')
```

Calculate the percentage of values being NaN for each column

```
In [ ]: # Calculate the percentage of values being NaN for each column
rows = len(df)
missing = df.isna().sum()
percentage_missing = missing / rows
```

```
In [ ]: # Put the data in a DataFrame and sort it
percentage_missing_df = pd.DataFrame({'Missing' : percentage_missing})
percentage_missing_df.sort_values('Missing', ascending = False, inplace = True)
```

```
In [ ]: #printing columns with more than 10% missing values
print(percentages_missing_df[percentages_missing_df['Missing']>0.1])
```

	Missing
Schedule	0.858453
Air.carrier	0.812710
FAR.Description	0.639742
Aircraft.Category	0.636772
Longitude	0.613304
Latitude	0.613203
Airport.Code	0.436016
Airport.Name	0.407081
Broad.phase.of.flight	0.305606
Publication.Date	0.154924
Total.Serious.Injuries	0.140737
Total.Minor.Injuries	0.134246
Total.Fatal.Injuries	0.128261

## Drop columns with over 50% missing values

```
In [ ]: # Drop columns with over 50% missing values
cols_to_drop = list(percentages_missing_df[percentages_missing_df['Missing'] > 0.5].i
df.drop(columns = cols_to_drop, axis = 1, inplace = True)
print(cols_to_drop)
```

```
['Schedule', 'Air.carrier', 'FAR.Description', 'Aircraft.Category', 'Longitude', 'La
titude']
```

## Drop records not Accidents in United States

```
In [ ]: # Drop records not Accidents in United States
before = len(df)
df = df[(df['Investigation.Type'] == 'Accident') & (df['Country'] == 'United States')]
dropped = before - len(df)
print(str(dropped) + ' rows dropped.')
```

8983 rows dropped.

```
In [ ]: # Convert Date to a datetime, add a Year & Month column and remove data before 1982
df['Event.Date'] = pd.to_datetime(df['Event.Date'])
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 79906 entries, 0 to 88888
Data columns (total 25 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Event.Id                             79906 non-null  object
1   Investigation.Type                    79906 non-null  object
2   Accident.Number                      79906 non-null  object
3   Event.Date                           79906 non-null  datetime64[ns]
4   Location                             79895 non-null  object
5   Country                              79906 non-null  object
6   Airport.Code                         47449 non-null  object
7   Airport.Name                        49918 non-null  object
8   Injury.Severity                     79854 non-null  object
9   Aircraft.damage                     78782 non-null  object
10  Registration.Number                 79838 non-null  object
11  Make                               79894 non-null  object
12  Model                              79877 non-null  object
13  Amateur.Built                      79891 non-null  object
14  Number.of.Engines                  78147 non-null  float64
15  Engine.Type                        76988 non-null  object
16  Purpose.of.flight                 78025 non-null  object
17  Total.Fatal.Injuries               69641 non-null  float64
18  Total.Serious.Injuries             68921 non-null  float64
19  Total.Minor.Injuries               69551 non-null  float64
20  Total.Uninjured                   74911 non-null  float64
21  Weather.Condition                 79345 non-null  object
22  Broad.phase.of.flight              59297 non-null  object
23  Report.Status                     77341 non-null  object
24  Publication.Date                   67649 non-null  object
dtypes: datetime64[ns](1), float64(5), object(19)
memory usage: 15.9+ MB
```

```
In [ ]: #Add a day, month & year column
df['Year'] = df['Event.Date'].dt.year
df['Month.Abbbr'] = df['Event.Date'].dt.month_name().str[:3]
df['Day.Name.Abbbr'] = df['Event.Date'].dt.day_name().str[:3]

# Add a weekend column
df.loc[(df['Day.Name.Abbbr'] == 'Sat') | (df['Day.Name.Abbbr'] == 'Sun'), 'Weekend'] = 1
df.loc[(df['Day.Name.Abbbr'] != 'Sat') & (df['Day.Name.Abbbr'] != 'Sun'), 'Weekend'] = 0

# Remove data before 1982
df = df[df['Year'] >= 1982]
```

```
In [ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 79899 entries, 7 to 88888
Data columns (total 29 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Event.Id                             79899 non-null  object
1   Investigation.Type                   79899 non-null  object
2   Accident.Number                     79899 non-null  object
3   Event.Date                          79899 non-null  datetime64[ns]
4   Location                            79888 non-null  object
5   Country                             79899 non-null  object
6   Airport.Code                        47449 non-null  object
7   Airport.Name                        49918 non-null  object
8   Injury.Severity                     79847 non-null  object
9   Aircraft.damage                     78775 non-null  object
10  Registration.Number                 79831 non-null  object
11  Make                               79887 non-null  object
12  Model                              79870 non-null  object
13  Amateur.Built                      79884 non-null  object
14  Number.of.Engines                  78141 non-null  float64
15  Engine.Type                        76982 non-null  object
16  Purpose.of.flight                 78019 non-null  object
17  Total.Fatal.Injuries               69635 non-null  float64
18  Total.Serious.Injuries             68916 non-null  float64
19  Total.Minor.Injuries               69546 non-null  float64
20  Total.Uninjured                    74905 non-null  float64
21  Weather.Condition                  79338 non-null  object
22  Broad.phase.of.flight              59290 non-null  object
23  Report.Status                      77334 non-null  object
24  Publication.Date                   67643 non-null  object
25  Year                               79899 non-null  int32
26  Month.Abbbr                       79899 non-null  object
27  Day.Name.Abbbr                    79899 non-null  object
28  Weekend                           79899 non-null  object
dtypes: datetime64[ns](1), float64(5), int32(1), object(22)
memory usage: 18.0+ MB
```

```
In [ ]: # Merge same airport names together
df['Airport.Name'].replace(to_replace = '(?i)^.*private.*$', value = 'PRIVATE', inplace = True)
df['Airport.Name'].replace(to_replace = '(?i)none', value = 'NONE', inplace = True)
df['Airport.Name'].value_counts().nlargest(10)
```

C:\Users\Admin\AppData\Local\Temp\ipykernel\_5972\2540906905.py:2: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.  
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df['Airport.Name'].replace(to_replace = '(?i)^.*private.*$', value = 'PRIVATE', inplace = True, regex = True)
```

C:\Users\Admin\AppData\Local\Temp\ipykernel\_5972\2540906905.py:3: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df['Airport.Name'].replace(to_replace = '(?i)none', value = 'NONE', inplace = True, regex = True)
```

```
Out[ ]: Airport.Name
PRIVATE      1204
NONE          143
MERRILL FIELD    83
VAN NUYS        79
MUNICIPAL       75
CENTENNIAL      74
UNKNOWN         68
CHINO           53
BIRCHWOOD       49
SEDONA          47
Name: count, dtype: int64
```

```
In [ ]: # Merge same registration numbers together
df['Registration.Number'].replace(to_replace = '(?i)none', value = 'NONE', inplace = True)
df['Registration.Number'].value_counts().nlargest(10)
```

C:\Users\Admin\AppData\Local\Temp\ipykernel\_5972\2380449474.py:2: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df['Registration.Number'].replace(to_replace = '(?i)none', value = 'NONE', inplace = True, regex = True)
```

```
Out[ ]: Registration.Number
      NONE      341
      UNREG     114
      N20752      7
      N4101E      6
      N11VH      6
      N8402K      6
      N53893      6
      N5408Y      6
      N121CC      6
      N3331R      5
      Name: count, dtype: int64
```

```
In [ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 79899 entries, 7 to 88888
Data columns (total 29 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Event.Id                             79899 non-null  object
1   Investigation.Type                    79899 non-null  object
2   Accident.Number                      79899 non-null  object
3   Event.Date                           79899 non-null  datetime64[ns]
4   Location                             79888 non-null  object
5   Country                              79899 non-null  object
6   Airport.Code                         47449 non-null  object
7   Airport.Name                         49918 non-null  object
8   Injury.Severity                     79847 non-null  object
9   Aircraft.damage                     78775 non-null  object
10  Registration.Number                 79831 non-null  object
11  Make                               79887 non-null  object
12  Model                              79870 non-null  object
13  Amateur.Built                      79884 non-null  object
14  Number.of.Engines                  78141 non-null  float64
15  Engine.Type                       76982 non-null  object
16  Purpose.of.flight                 78019 non-null  object
17  Total.Fatal.Injuries              69635 non-null  float64
18  Total.Serious.Injuries            68916 non-null  float64
19  Total.Minor.Injuries              69546 non-null  float64
20  Total.Uninjured                   74905 non-null  float64
21  Weather.Condition                 79338 non-null  object
22  Broad.phase.of.flight             59290 non-null  object
23  Report.Status                     77334 non-null  object
24  Publication.Date                  67643 non-null  object
25  Year                              79899 non-null  int32
26  Month.Abbbr                      79899 non-null  object
27  Day.Name.Abbrr                   79899 non-null  object
28  Weekend                           79899 non-null  object
dtypes: datetime64[ns](1), float64(5), int32(1), object(22)
memory usage: 18.0+ MB
```

## Merge different capitalizations of Make together

```
In [ ]: # Merge different capitalizations of Make together
df['Make'] = df['Make'].str.title()
df['Make'].value_counts().nlargest(10)
```



```
Out[ ]: Make
      Cessna      25566
      Piper      14008
      Beech       4892
      Bell       2236
      Mooney      1272
      Grumman     1131
      Bellanca    1036
      Boeing      931
      Robinson    916
      Hughes      868
      Name: count, dtype: int64
```

```
In [ ]: # Transform Amateur Built to boolean
df['Amateur.Built'].replace(to_replace = ['Yes', 'Y'], value = True, inplace = True)
df['Amateur.Built'].replace(to_replace = ['No', 'N'], value = False, inplace = True)
df['Amateur.Built'].value_counts()
```

C:\Users\Admin\AppData\Local\Temp\ipykernel\_5972\3362052493.py:2: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df['Amateur.Built'].replace(to_replace = ['Yes', 'Y'], value = True, inplace = True, regex = False)
```

```
Out[ ]: Amateur.Built
      False      71589
      True       8295
      Name: count, dtype: int64
```

## Split location in city and state

```
In [ ]: # Split Location in city and state
df['City'] = df['Location'].str.split(',').str[0]
df['State'] = df['Location'].str.split(',').str[1]
df[['City', 'State']].head(10)
```

Out[ ]:

	City	State
7	PULLMAN	WA
8	EAST HANOVER	NJ
9	JACKSONVILLE	FL
10	HOBBS	NM
11	TUSKEGEE	AL
12	HOMER	LA
13	HEARNE	TX
14	CHICKASHA	OK
15	LITTLE ROCK	AR
16	MIDWAY	UT

## Categorize the amount of injuries as this is already in another column

```
In [ ]: # Remove amount of injuries as this is already in another column
df['Injury.Severity'] = df['Injury.Severity'].str.split('(').str[0]
df['Injury.Severity'].value_counts()
```

```
Out[ ]: Injury.Severity
Non-Fatal      64457
Fatal          15019
Minor           203
Serious         153
Unavailable     15
Name: count, dtype: int64
```

```
In [ ]: # Merge weather condition unknowns
df['Weather.Condition'].replace(to_replace = ['Unk', 'UNK'], value = 'Unknown', inplace=True)
df['Weather.Condition'].value_counts()
```

C:\Users\Admin\AppData\Local\Temp\ipykernel\_5972\1600600250.py:2: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or 'df[col] = df[col].method(value)' instead, to perform the operation inplace on the original object.

```
df['Weather.Condition'].replace(to_replace = ['Unk', 'UNK'], value = 'Unknown', inplace = True, regex = False)
```

```
Out[ ]: Weather.Condition
      VMC      73340
      IMC      5387
      Unknown   611
      Name: count, dtype: int64
```

```
In [ ]: df.shape
```

```
Out[ ]: (79899, 31)
```

```
In [ ]: df.columns
```

```
Out[ ]: Index(['Event.Id', 'Investigation.Type', 'Accident.Number', 'Event.Date',
              'Location', 'Country', 'Airport.Code', 'Airport.Name',
              'Injury.Severity', 'Aircraft.damage', 'Registration.Number', 'Make',
              'Model', 'Amateur.Built', 'Number.ofEngines', 'Engine.Type',
              'Purpose.of.flight', 'Total.Fatal.Injuries', 'Total.Serious.Injuries',
              'Total.Minor.Injuries', 'Total.Uninjured', 'Weather.Condition',
              'Broad.phase.of.flight', 'Report.Status', 'Publication.Date', 'Year',
              'Month.Abb', 'Day.Name.Abb', 'Weekend', 'City', 'State'],
              dtype='object')
```

```
In [ ]: df.head()
```

Out[ ]:

	Event.Id	Investigation.Type	Accident.Number	Event.Date	Location	Cou
7	20020909X01562	Accident	SEA82DA022	1982-01-01	PULLMAN, WA	Un St
8	20020909X01561	Accident	NYC82DA015	1982-01-01	EAST HANOVER, NJ	Un St
9	20020909X01560	Accident	MIA82DA029	1982-01-01	JACKSONVILLE, FL	Un St
10	20020909X01559	Accident	FTW82DA034	1982-01-01	HOBBS, NM	Un St
11	20020909X01558	Accident	ATL82DKJ10	1982-01-01	TUSKEGEE, AL	Un St

5 rows × 31 columns



```
In [ ]: injury_data = df[df['Injury.Severity'] != 'Unavailable']
```

## Number of accidents per year


```
In [ ]: # Number of accidents per year
accidents_per_year = df.groupby(['Year'], as_index = False)['Event.Id'].count()
accidents_per_year
```

Out[ ]:

	Year	Event.Id
0	1982	3445
1	1983	3391
2	1984	3290
3	1985	2972
4	1986	2758
5	1987	2665
6	1988	2537
7	1989	2410
8	1990	2403
9	1991	2350
10	1992	2233
11	1993	2173
12	1994	2135
13	1995	2153
14	1996	2048
15	1997	1970
16	1998	2032
17	1999	2030
18	2000	1982
19	2001	1857
20	2002	1833
21	2003	1891
22	2004	1737
23	2005	1804
24	2006	1613
25	2007	1763
26	2008	1642
27	2009	1549
28	2010	1496
29	2011	1551

	Year	Event.Id
30	2012	1533
31	2013	1282
32	2014	1281
33	2015	1289
34	2016	1317
35	2017	1311
36	2018	1336
37	2019	1296
38	2020	1131
39	2021	1173
40	2022	1237

```
In [ ]: plot = sns.lineplot(x = 'Year', y = 'Event.Id', data = accidents_per_year, color =
```



## Number of fatal accidents per year

```
In [ ]: # Number of fatal accidents per year
fatal_accidents_per_year = df[df['Injury.Severity'] == 'Fatal'].groupby(['Year'], a
```


fatal\_accidents\_per\_year

Out[ ]:

	Year	Event.Id
0	1982	644
1	1983	654
2	1984	617
3	1985	558
4	1986	516
5	1987	489
6	1988	492
7	1989	479
8	1990	499
9	1991	478
10	1992	484
11	1993	426
12	1994	438
13	1995	430
14	1996	381
15	1997	371
16	1998	383
17	1999	352
18	2000	373
19	2001	351
20	2002	363
21	2003	377
22	2004	349
23	2005	335
24	2006	322
25	2007	303
26	2008	283
27	2009	268
28	2010	265
29	2011	278

	Year	Event.Id
30	2012	270
31	2013	220
32	2014	253
33	2015	235
34	2016	210
35	2017	205
36	2018	224
37	2019	240
38	2020	200
39	2021	200
40	2022	204

```
In [ ]: sns.lineplot(x = 'Year', y = 'Event.Id', data = fatal_accidents_per_year, color = 'blue')
Out[ ]: <Axes: xlabel='Year', ylabel='Event.Id'>
```



```
In [ ]: # Calculate average fatality rate
averagefatal = len(injury_data[injury_data['Injury.Severity'] == 'Fatal'].index) /
print("Average fatality rate: " + str(round(averagefatal * 100, 2)) + '%')
```




Average fatality rate: 18.8%

## Months with the most accidents

```
In [ ]: # Months with the most accidents
plot = sns.countplot(x = 'Month.Abbbr', color = 'b', data = df)
plot.set(xlabel = 'Month', ylabel = None, title = 'Number of aviation accidents per
```


```
Out[ ]: [Text(0.5, 0, 'Month'),
Text(0, 0.5, ''),
Text(0.5, 1.0, 'Number of aviation accidents per month')]
```



## Days with the most accidents

```
In [ ]: # Days with the most accidents
plot = sns.countplot(x = 'Day.Name.Abbbr', order = ['Mon', 'Tue', 'Wed', 'Thu', 'Fri', 'S
plot.set(xlabel = 'day', ylabel = None, title = 'Number of aviation accidents per d
```

```
Out[ ]: [Text(0.5, 0, 'day'),
Text(0, 0.5, ''),
Text(0.5, 1.0, 'Number of aviation accidents per day of the week')]
```



## Number of fatal accidents per year against the model.

With accidents per year against the model will help us know which model caused more accidents at a particular year.


```
In [ ]: #Number of fatal accidents per year against the model  
accidents_per_years = df.groupby(['Year'], as_index = False)['Model'].count()  
accidents_per_years
```

Out[ ]:

	Year	Model
0	1982	3431
1	1983	3391
2	1984	3289
3	1985	2972
4	1986	2757
5	1987	2665
6	1988	2537
7	1989	2409
8	1990	2403
9	1991	2350
10	1992	2230
11	1993	2173
12	1994	2134
13	1995	2153
14	1996	2048
15	1997	1970
16	1998	2032
17	1999	2029
18	2000	1981
19	2001	1857
20	2002	1833
21	2003	1889
22	2004	1736
23	2005	1804
24	2006	1611
25	2007	1763
26	2008	1642
27	2009	1549
28	2010	1496
29	2011	1551

	Year	Model
30	2012	1533
31	2013	1282
32	2014	1280
33	2015	1289
34	2016	1317
35	2017	1311
36	2018	1336
37	2019	1296
38	2020	1131
39	2021	1173
40	2022	1237

```
In [ ]: plot = sns.lineplot(x = 'Year', y = 'Model', data = accidents_per_years, color = '#
```



```
In [ ]: df
```

Out[ ]:

Event.IdInvestigation.TypeAccident.NumberEvent.DateLocation

7	20020909X01562	Accident	SEA82DA022	1982-01-01	PULLMAN, WA
8	20020909X01561	Accident	NYC82DA015	1982-01-01	EAST HANOVER, NJ
9	20020909X01560	Accident	MIA82DA029	1982-01-01	JACKSONVILLE, FL
10	20020909X01559	Accident	FTW82DA034	1982-01-01	HOBBS, NM
11	20020909X01558	Accident	ATL82DKJ10	1982-01-01	TUSKEGEE, AL
...	...	...	...	...	...
88884	20221227106491	Accident	ERA23LA093	2022-12-26	Annapolis, MD
88885	20221227106494	Accident	ERA23LA095	2022-12-26	Hampton, NH
88886	20221227106497	Accident	WPR23LA075	2022-12-26	Payson, AZ
88887	20221227106498	Accident	WPR23LA076	2022-12-26	Morgan, UT
88888	20221230106513	Accident	ERA23LA097	2022-12-29	Athens, GA

79899 rows × 31 columns



In [ ]:

df.to\_csv("New\_Cleaned\_Data.csv", index = False )